

Command Line Interface, Stochastic SVD*

Contents

1	Background information	2
1.1	Stochastic SVD (SSVD)	2
1.2	PCA options in SSVD	4
2	Input/output file formats and layout	7
3	CLI usage	7
4	Embedded use	10
5	Mixed environments: R+Mahout	10

*Dmitriy Lyubimov, dlyubimov at apache dot org

1 Background information

1.1 Stochastic¹ SVD (SSVD)

Stochastic SVD method in Mahout produces reduced rank Singular Value Decomposition output in its strict mathematical definition:

$$\mathbf{A} \approx \mathbf{U}\mathbf{\Sigma}\mathbf{V}^\top,$$

i. e. it creates outputs for matrices \mathbf{U} , \mathbf{V} and $\mathbf{\Sigma}$, each of which may be requested individually. The desired rank of decomposition, henceforth denoted as $k \in \mathbb{N}_1$, is a parameter of the algorithm. The singular values inside diagonal matrix $\mathbf{\Sigma}$ satisfy $\sigma_{i+1} \leq \sigma_i \forall i \in [1, k-1]$, i.e. sorted from biggest to smallest. Cases of rank deficiency $\text{rank}(\mathbf{A}) < k$ are handled by producing 0s in singular value positions once deficiency takes place.

Single space for comparing row-items and column-items. On top of it, there's an option to present decomposition output in a form of

$$\mathbf{A} \approx \mathbf{U}_\sigma \mathbf{V}_\sigma^\top, \tag{1}$$

where one can request $\mathbf{U}_\sigma = \mathbf{U}\mathbf{\Sigma}^{0.5}$ instead of \mathbf{U} (but not both), $\mathbf{V}_\sigma = \mathbf{V}\mathbf{\Sigma}^{0.5}$ instead of \mathbf{V} (but not both). Here, notation $\mathbf{\Sigma}^{0.5}$ implies diagonal matrix containing square roots of the singular values:

$$\mathbf{\Sigma}^{0.5} = \begin{pmatrix} \sqrt{\sigma_1} & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & \sqrt{\sigma_k} \end{pmatrix}.$$

Original singular values $\mathbf{\Sigma}$ are still produced and saved regardless.

This option is a nod to a common need of comparing actors represented by both input rows and input columns in a common space. E.g. if LSI is performed such that rows are documents and columns are terms then it is possible to compare documents and terms (either existing or fold in new ones) in one common space and perform similarity measurement between a document and a term, rather than computing just term2term or document2document similarities.

¹See "A note about stochasticity and result precision."

Folding in new observations. It is probably worth mentioning the operation of “folding in” new observations in context of this method, since it is often a basis for incremental methods.

If $\tilde{\mathbf{c}}_r$ ($\tilde{\mathbf{c}}_c$) is a new row (column) observation in addition to original input \mathbf{A} , then correspondent “new” row vectors of $\tilde{\mathbf{U}}$ ($\tilde{\mathbf{V}}$) can be obtained as

$$\tilde{\mathbf{u}} = \Sigma^{-1} \mathbf{V}^\top \tilde{\mathbf{c}}_r, \quad (2)$$

$$\tilde{\mathbf{v}} = \Sigma^{-1} \mathbf{U}^\top \tilde{\mathbf{c}}_c. \quad (3)$$

Similarly, for the case (1) folding in new observations into rows of $\tilde{\mathbf{U}}_\sigma$ ($\tilde{\mathbf{V}}_\sigma$) would look like

$$\tilde{\mathbf{u}}_\sigma = \mathbf{V}_\sigma^\top \tilde{\mathbf{c}}_r, \quad (4)$$

$$\tilde{\mathbf{v}}_\sigma = \mathbf{U}_\sigma^\top \tilde{\mathbf{c}}_c. \quad (5)$$

Thus, new rows can be added to matrices denoted as $\tilde{\mathbf{U}}$ ($\tilde{\mathbf{V}}$) corresponding to new observations as new observations become available, i.e. incrementally. Given that new observations are usually moderately sparse vectors, it might be feasible to do fold-in in real time or almost real time, assuming proper fast row-wise indexing of \mathbf{U} (\mathbf{V}) exists (e.g. using a batch request to an HBase table containing rows of \mathbf{U} (\mathbf{V})). However, since operation of folding in new observations doesn’t change original decomposition and its spaces, such new observations cannot be considered ‘training’ examples. Typically, from time to time accumulated new observations can be added to original input \mathbf{A} and the whole decomposition can be recomputed again.

Common applications for SVD include Latent Semantic Analysis (LSA), Principal Component Analysis (PCA), dimensionality reduction and others.

A note about stochasticity and result precision. The “stochastic” term in this case has somewhat different implications from a common expectation we came to have by seeing this epithet next to a computational method name. Unlike some other stochastic approaches (such as stochastic gradient descent) this method is not affected by effect of using random numbers quite directly. Randomicity and probability is only used in the first phase to find a suitable reduced subspace of original space. But once the subspace is selected, the problem is projected into that space analytically and then the whole problem is solved (and rotated out of the solution subspace) 100% analytically, so there’s no requirement per se for such initial subspace to be found with a super great accuracy but only with some acceptable probability of an acceptable error w.r.t degenerate projection cases. From here on, all obtained errors in the final solution are entirely due to the fact that some “important” data points were actually projected onto subspace in a degenerate way resulting in high relative rounding errors due to limits posed by machine precision. There’s also a subsequent non-stochastic step known as “power iterations” to improve initial probabilistic guess on such subspace and of

what “data points” are really important to project in non-degenerate way which, as the real-life experimental data has shown, all but eliminates any further practical precision concerns even in comparison to the Lanczos approach in Mahout².

Because of the said, I came to think to myself of this approach as “semi-stochastic” (and “semi-analytical”) to emphasize the fact that all final errors are actually rounding errors (and their probabilistic estimate also depends on machine precision available perhaps more prominently than in most other stochastic approaches). In other words, stochasticity affects *stability* of the ad-hoc solution only. Strictly speaking, if we could reduce machine precision error to 0, the method would have always been giving optimal solution regardless of effects of stochasticity (due to the fact that strictly speaking a probability of multiplication of a non-zero vector by a random vector being exactly 0 (or exactly anything, for this matter) is also 0). This is generally not true with most other stochastic methods in a sense that reducing machine precision to 0 per se will not result in error estimate also being 0 if compared to the optimal solution.

With respect to all the said, the case that needs to be specifically warned about is the case where the data has variances in a *great deal* of principal orthogonal directions *with no statistically significant differences* in between them, i.e., the input data has *no interesting spectrum decay*. A matrix where each entry is generated from a random distribution with 0 mean, would be an example of such input data. Figuring principal data variances for such randomly generated and at the same time “big” data are hard (in part due to effects of law of large numbers) and may result in sufficiently problematic rounding errors in the final solution compared to optimal solution (if we could obtain it at that scale). But the cases dominated by a random noise with no statistically significant signal are hopefully not the data we actually are forced to face with in practice. But even quite small signal buried in a sea of “big data” noise could probably eventually be dug out given enough computational effort with spent, with this method (by cranking up the `--q` setting) but will likely also eventually reach a point of diminishing return.

1.2 PCA options in SSVD

Some of interesting applications of SVD is dimensionality reduction and Principal Component Analysis. As of MAHOUT-817, SSVD method is equipped with options helping to produce both PCA and dimensionality reduction transformations.

Outline: data points are row vectors. We approach general PCA and dimensionality reduction problem with respect to input expressed in Mahout’s distributed row matrix format. We also assume data points are row vectors in such matrix³. We denote such $m \times n$ input matrix as $\mathbf{A}^{(pca)}$:

$$\mathbf{A}^{(pca)} = \begin{pmatrix} \mathbf{a}_1^\top \\ \mathbf{a}_2^\top \\ \vdots \\ \mathbf{a}_m^\top \end{pmatrix}$$

²TODO: get permission and add reference to Nathan Halko’s work w.r.t. precision comparisons

³Note that in wikipedia article input data points are considered to be column vectors, so our input is transpose w.r.t. to case outlined in the wikipedia PCA article.

Column mean is n-vector⁴

$$\begin{aligned}\xi &= \begin{pmatrix} \xi_1 \\ \xi_2 \\ \vdots \\ \xi_n \end{pmatrix} \\ &= \frac{1}{m} \sum_i^m \mathbf{a}_i.\end{aligned}$$

We denote $m \times n$ mean matrix as

$$\Xi = \begin{pmatrix} \xi^\top \\ \xi^\top \\ \vdots \\ \xi^\top \end{pmatrix} \in \mathbb{R}^{m \times n}$$

Traditional approach under these settings starts with finding a column mean ξ and subtract it from all data points (row vectors) of the input

$$\mathbf{A} = \mathbf{A}^{(pca)} - \Xi$$

and then proceeds with finding a reduced k -rank SVD:

$$\mathbf{A} \approx \mathbf{U} \Sigma \mathbf{V}^\top.$$

At this point rows of \mathbf{U} correspond to original data points (rows of $\mathbf{A}^{(pca)}$) converted to PCA space and \mathbf{V} represents approximate eigenvectors of the covariance matrix of our input (and we are done with PCA part at this point). Note that since our datapoints in this case are row vectors in the input (and not column vectors), covariance matrix is $\mathbf{C} = \frac{1}{m} \mathbf{A}^\top \mathbf{A}$.⁵

⁴In many cases in literature one would find PCA mean vector denoted as μ .

⁵Some courses (see Andrew Ng's ML online lectures) also mention a metric aka "percent of variance retained" $\sum_{i=1}^k \sigma_i / \sum_{i=1}^n \sigma_i \cdot 100\%$ that seems to suggest that it would be a metric to compare "goodness" of choice for k . We of course cannot know all of the singular values of the full SVD in context of SSVD, although we can estimate that

$$\frac{\sum_{i=1}^k \sigma_i}{\sum_{i=1}^n \sigma_i} \geq \frac{\sum_{i=1}^k \sigma_i}{(n-k) \sigma_k + \sum_{i=1}^k \sigma_i},$$

i.e. we can make a statement in a sense that "we have *at least* that much variance retained" instead. Whether this estimate is useful or not in practice, is not clear to me as I have no error estimate (and, more importantly, no error estimate *with* stochastic error baked in) between expression on left and right of the inequality at the moment. But at least it puts some metric threshold on how big k we may need and identify some cases where smaller values of k will suffice given particular input and minimum variance retained prerequisite thus reducing need for flops. Another possible way to estimate it is perhaps to try and fit existing singular values into asymptotically decaying curve in order to produce a better guess for $\sum_{i=k+1}^n \sigma_i$.

Transformations to/from for new observation data points. Dimensionality reduction transformations are directly following from SVD fold-in operation (2) described above.

Transformation of any new data point observation of an n -vector $\tilde{\mathbf{c}}_r$ into reduced dimensionality PCA space k -vector $\tilde{\mathbf{u}}$ will look like

$$\tilde{\mathbf{u}} = \mathbf{\Sigma}^{-1} \mathbf{V}^\top (\tilde{\mathbf{c}}_r - \mathbf{\xi}) \quad (6)$$

(this operation is essentially an SVD fold-in operation corrected for the mean subtraction). Note that, again, if input vectors tend to be quite sparse, then (6) could be decomposed as

$$\tilde{\mathbf{u}} = \mathbf{\Sigma}^{-1} \mathbf{V}^\top \tilde{\mathbf{c}}_r - \mathbf{\Sigma}^{-1} \mathbf{V}^\top \mathbf{\xi},$$

and online conversion can be sped up by precomputing term $\mathbf{\Sigma}^{-1} \mathbf{V}^\top \mathbf{\xi}$ which ends up to be a small dense k -vector, and big matrix \mathbf{V} could be row-indexed for fast online matrix-vector multiplication.

Inverse transformation (from reduced PCA space into original space) looks like

$$\tilde{\mathbf{c}}_r = \mathbf{\xi} + \mathbf{V} \mathbf{\Sigma} \tilde{\mathbf{u}}. \quad (7)$$

MAHOUT-817 goals: why brute force approach is hard in context of big data computations. In context of massive computations, input $\mathbf{A}^{(pca)}$ is often rather sparse. Sparse matrices are packed and their subsequent operations within Mahout framework are optimized to account for degenerate nature of zero elements computation. Sometimes such reduction of need for flops and space may approach several orders of magnitude. However, mean subtraction step would turn such sparse inputs into a dense matrix $(\mathbf{A}^{(pca)} - \mathbf{\Xi})$. Such intermediate input will take a lot of space and subsequent SVD will require a lot of flops. Fortunately, this can be addressed in context of SSVD method in a way that will be (almost) cost-equivalent to a regular SSVD computation on original sparse input $\mathbf{A}^{(pca)}$.

MAHOUT-817 addresses two goals:

- Provide column-wise mean computation step in the whole pipeline (or use outside mean vector if already available)
- Lift the dense matrix data concerns per above.

The sparser the original input is, the more efficiency gain is to be had by using SSVD PCA options compared to brute-force approach.

If original input is 100% dense, SSVD PCA options will have roughly the same cost as brute-force approach.

MAHOUT-817 introduces two additional pca options: `--pca` and `--pcaOffset` that request to treat incoming data as a PCA input. SVD rank option `-k` will correspond to the reduced dimensionality of PCA space. See §3 for details.

2 Input/output file formats and layout

Input \mathbf{A} , as well as outputs \mathbf{U} (\mathbf{U}_σ), \mathbf{V} (\mathbf{V}_σ), are in Mahout's Distributed Row Matrix format, i.e. set of sequence files where value is of `VectorWritable` type. As far as keys are concerned, rows of \mathbf{A} may be keyed (identified) by any `Writable` (for as long as it is instantiable thru a default constructor). That, among other things, means that this method can be applied directly on the output of `seq2sparse` where keys are of `Text` type⁶.

Definition of output \mathbf{U} (\mathbf{U}_σ) is identical to definition of the input matrix \mathbf{A} , and the keys of corresponding rows in \mathbf{A} are copied to corresponding rows of output \mathbf{U} (\mathbf{U}_σ).

Definition of output \mathbf{V} (\mathbf{V}_σ) is always sequence file(s) of (`IntWritable`, `VectorWritable`) where key corresponds to a column index of the input \mathbf{A} .

Output of Σ is encoded by a single output file with a single vector value (`VectorWritable`) with main diagonal entries of Σ aka singular values ($\sigma_1 \dots \sigma_k$).

3 CLI usage⁷

```
mahout ssvd <options>
```

⁶(TODO: re-verify)

⁷As of Mahout 0.7 trunk.

12/13/2011 adjusted for MAHOUT-922 changes.

02/22/2012 MAHOUT-817.

Options.

- k, --rank <int-value> (required): the requested SVD rank (minimum number of singular values and dimensions in U, V matrices). The value of $k + p$ directly impacts running time and memory requirements. ***$k+p=500$ is probably more than reasonable.*** Typically $k + p$ is taken within range 20...200.
- p, --oversampling <int-value> (optional, default 15): stochastic SVD oversampling. p doesn't seem to have to be very significant. If power iterations ($q > 0$) are used then p perhaps could be kept quite low, not to exceed 10% of k .
- q, --powerIter <int-value> (optional, default 0): number of power iterations to perform. This helps fighting data noise and improve precision significantly more than just increasing p . Each additional power iteration adds 2 more steps (map/reduce + map-only). Experimental data suggests using $q = 1$ is already producing quite good results which are hard to much improve upon.
- r, --blockHeight <int-value> (optional, default 10,000): the number of rows of source matrix for block computations during $\mathbf{Y} = \mathbf{Q}\mathbf{R}$ decomposition. Taller blocking causes more memory use but produces less blocks and therefore somewhat better running times. The most optimal mode from the running time point of view should be 1 block per 1 mapper. *This cannot be less than $k+p$.*
- oh, --outerProdBlockHeight <int-value> (optional, default 30,000): the block height during $\mathbf{B} = \mathbf{Q}^\top \mathbf{A}$ operation⁸⁹.
- abth, --abtBlockHeight <int-value> (optional, default 200,000): the block height during $\mathbf{Y}_i = \mathbf{A}\mathbf{B}_i^\top$ multiplication⁸⁹.
- s, --minSplitSize <int-value> (optional, default: use Hadoop's default): minimum split size to use in mappers reading \mathbf{A} input.¹⁰

⁸With extreme sparse matrices increasing this parameter may lead to better performance by reducing computational pressure on the shuffle and sort and grouping sparse records together.

⁹ Watch for GC thrashing and swap. Values too high may cause GC thrashing and/or swapping, both of which are capable of bringing job performance down to a halt. Don't starve jobs for memory. Defaults are believed to work well for -Xmx800mb or above per child task.

¹⁰*As of this day, I haven't heard of a case where somebody would actually have to use this option and actually increase split size and how it has played out. So this option is experimental.*

Since in this version projection block formation happens in mappers, for a sufficiently wide input matrix the algorithm may not be able to read minimum $k + p$ rows and form a block of minimum height required, so in that case the job would bail out at the very first mapping step. If this happens, one of the recourses available is to force increase in the MapReduce split size using `SequenceFileInputFormat.setMinSplitSize()` property. Increasing this significantly over HDFS block size may result in network IO to mappers. Another caveat is that one sometimes

--computeU <true|false> (optional, default true). Request computation of the \mathbf{U} matrix

--computeV <true|false> (optional, default true). Request computation of the \mathbf{V} matrix

--vHalfSigma <true|false> (optional, default: false): compute $\mathbf{V}_\sigma = \mathbf{V}\Sigma^{0.5}$ instead of \mathbf{V} (see overview for explanation).

--uHalfSigma <true|false> (optional, default: false): compute $\mathbf{U}_\sigma = \mathbf{U}\Sigma^{0.5}$ instead of \mathbf{U} .

--reduceTasks <int-value> optional. The number of reducers to use (where applicable): depends on the size of the hadoop cluster. At this point it could also be overwritten by a standard hadoop property using -D option⁶. ***Probably always needs to be specified as by default Hadoop would set it to 1, which is certainly far below the cluster capacity.*** Recommended value for this option $\sim 95\%$ or $\sim 190\%$ of available reducer capacity to allow for opportunistic executions.

--pca run in pca mode: treat the input as $\mathbf{A}^{(pca)}$ and also compute column-wise mean ξ over the input and use it to compute PCA space (unless external column mean is already provided by --pcaOffset). (see §1.2)

--pcaOffset <xi-path> (optional, default: none). Path(glob) of external column mean. The glob parameter must point to a sequence file containing single **VectorWritable** row as the ξ mean (see §1.2). This option can be used if column-wise mean is already efficiently obtained as byproduct from another pipeline, or if one wants to use custom centering offset for the data. This will save one MR pass over input since the mean will not have to be computed.

Standard Mahout options.

--input <glob> HDFS glob specification where the DistributedRowMatrix input to be found.

--output <hdfs-dir> non-existent hdfs directory where to output \mathbf{U} , \mathbf{V} and Σ (singular values) files.

does not want too many mappers because it may in fact increase time of the computation. Consequently, this option should probably be left alone unless one has significant amount of mappers (as in thousands of map tasks) at which point reducing amount of mappers may actually improve the throughput (just a guesstimate at this point).

`--tempDir <temp-dir>` temporary dir where to store intermediate files (cleaned up upon normal completion). This is a standard Mahout optional parameter.

`-ow, --overwrite` overwrite output if exists.

4 Embedded use

It is possible to instantiate and use `SSVDSolver` class in embedded fashion in Hadoop-enabled applications. This class would have getter and setter methods for each option available via command line. See javadoc for details.

5 Mixed environments: R+Mahout

wip, todo, check back later.